

CS420 Assignment 1

Question 1

1. $A \perp E \mid \{D\}$

- $A \rightarrow B \rightarrow D \rightarrow E$
 - Blocked due to cascade on D, it is observed.
- $A \rightarrow B \leftarrow C \rightarrow F \leftarrow E$
 - Blocked due to v-structure on F, neither F nor its descendants are observed.
- $A \rightarrow B \rightarrow D \rightarrow G \rightarrow H \leftarrow F \leftarrow E$
 - Blocked due to cascade on D, it is observed.
- $A \rightarrow D \rightarrow E$
 - Blocked due to cascade on D, it is observed.
- $A \rightarrow D \rightarrow G \rightarrow H \rightarrow F \rightarrow E$
 - Blocked due to cascade on D, it is observed.

Since all the trails are blocked, the statement is true.

2. $A \perp H \mid \{B, G\}$

- $A \rightarrow D \rightarrow G \rightarrow H$
 - Blocked due to cascade on D, it is observed.
- $A \rightarrow D \rightarrow E \rightarrow F \rightarrow H$
 - Blocked due to cascade on D, it is observed.
- $A \rightarrow D \leftarrow B \leftarrow C \rightarrow F \rightarrow H$
 - Active trail
- $A \rightarrow B \leftarrow C \rightarrow F \rightarrow H$
 - Active trail

- $A \rightarrow B \rightarrow D \rightarrow G \rightarrow H$
 - Blocked due to cascade on B, it is observed.
- $A \rightarrow B \rightarrow D \rightarrow E \rightarrow F \rightarrow H$
 - Blocked due to cascade on B, it is observed.
- $A \rightarrow B \leftarrow C \rightarrow F \leftarrow E \leftarrow D \rightarrow G \rightarrow H$
 - Blocked due to cascade on G, it is observed.

Since there are active trails, the statement is false.

3. $A \perp C$

- $A \rightarrow B \leftarrow C$
 - Blocked due to v-structure on B, neither B nor its descendants are observed.
- $A \rightarrow B \rightarrow D \rightarrow E \rightarrow F \leftarrow C$
 - Blocked due to v-structure on F, neither F nor its descendants are observed.
- $A \rightarrow B \rightarrow D \rightarrow G \rightarrow H \leftarrow F \leftarrow C$
 - Blocked due to v-structure on H, neither H nor its descendants are observed.
- $A \rightarrow D \rightarrow E \rightarrow F \leftarrow C$
 - Blocked due to v-structure on F, neither F nor its descendants are observed.
- $A \rightarrow D \leftarrow B \leftarrow C$
 - Blocked due to v-structure on D, neither D nor its descendants are observed.
- $A \rightarrow D \rightarrow G \rightarrow H \leftarrow F \leftarrow C$
 - Blocked due to v-structure on H, neither H nor its descendants are observed.

Since all the trails are blocked, the statement is true.


4. $\{D\} \perp \{C, F\} \mid \{A, B, E, G\}$

- $D \rightarrow E \rightarrow F \leftarrow C$
 - Blocked due to cascade on E, it is observed.
- $D \leftarrow B \leftarrow C \leftarrow F$
 - Blocked due to cascade on B, it is observed.
- $D \leftarrow A \rightarrow B \rightarrow C \rightarrow F$
 - Blocked due to common cause on A, it is observed.
- $D \rightarrow G \rightarrow H \rightarrow F \rightarrow C$
 - Blocked due to cascade on G, it is observed.

Since all the trails are blocked, the statement is true.

Question 2

A)



B)

Variable Name	Domain (list all the entries of the domain)
T	{Low actual temperature = 0, High actual temperature = 1}
G	{Low measured temperature = 0, High measured temperature = 1}
FG	{Gauge not faulty = 0, Gauge is faulty = 1}
FA	{Alarm not faulty = 0, Alarm is faulty = 1}
A	{Alarm does not sound = 0, Alarm sounds = 1}

C)

$$P(G = 1 \mid FG = 1, T) = 0.2$$

$$P(G = 0 \mid FG = 1, T) = 1 - 0.2 = 0.8$$

$$P(G = 1 \mid FG = 0, T) = 0.9$$

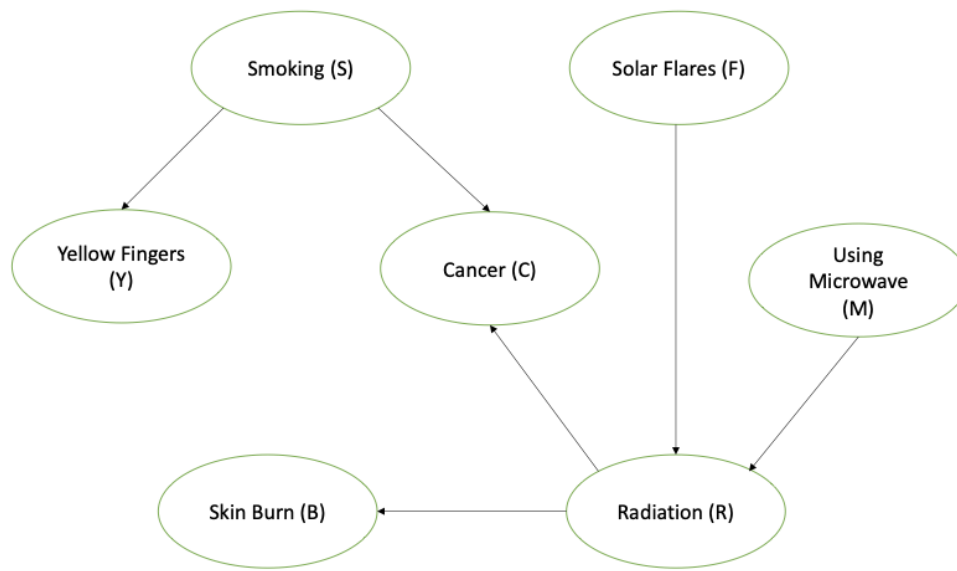
$$P(G = 0 \mid FG = 0, T) = 1 - 0.9 = 0.1$$

FG	T	$P(G = 1 \mid FG, T)$	$P(G = 0 \mid FG, T)$
1	1	0.2	0.8
1	0	0.8	0.2
0	1	0.9	0.1
0	0	0.1	0.9

D)

FA	G	$P(A = 1 \mid FA, G)$	$P(A = 0 \mid FA, G)$
0	0	1	0
0	1	1	0
1	0	0	1
1	1	0	1

Question 3



1.

```
[25] # Import relevant packages
from pgmpy.models import BayesianModel
from pgmpy.factors.discrete import TabularCPD
import sys

[26] # We first create a model which contains edges of the graph
model = BayesianModel([('Smoking', 'Yellow Fingers'), ('Smoking', 'Cancer'), ('Solar Flares', 'Radiation'), ('Using Microwave', 'Radiation'), ('Radiation',
# Enter conditional probability distribution for each variable

# Prior probability for Smoking P(S)
cpd_S = TabularCPD(variable='Smoking', variable_card=2, values=[[0.8], [0.2]])

# Prior probability for Solar Flares P(F)
cpd_F = TabularCPD(variable='Solar Flares', variable_card=2, values=[[0.999], [0.001]])

# Prior probability for Using Microwave P(M)
cpd_M = TabularCPD(variable='Using Microwave', variable_card=2, values=[[0.1], [0.9]])

WARNING:pgmpy:BayesianModel has been renamed to BayesianNetwork. Please use BayesianNetwork class, BayesianModel will be removed in future.
```

```
[27] # Conditional probability for Skin Burn or P(B|R)
cpd_B = TabularCPD(variable='Skin Burn', variable_card=2, values = [[0.99, 0.9], [0.01, 0.1]],
evidence = ['Radiation'],
evidence_card=[2])

# Conditional probability for Radiation or P(R|F, M)
cpd_R = TabularCPD(variable='Radiation', variable_card=2, values = [[0.99, 0.9, 0.7, 0.5], [0.01, 0.1, 0.3, 0.5]],
evidence = ['Solar Flares', 'Using Microwave'],
evidence_card=[2, 2])

# Conditional probability for Cancer or P(C|S, R)
cpd_C = TabularCPD(variable='Cancer', variable_card=2, values = [[0.9, 0.4, 0.8, 0.1], [0.1, 0.6, 0.2, 0.9]],
evidence = ['Smoking', 'Radiation'],
evidence_card=[2, 2])

# Conditional probability for Yellow Fingers or P(Y|S)
cpd_Y = TabularCPD(variable='Yellow Fingers', variable_card=2, values = [[0.9, 0.1], [0.1, 0.9]],
evidence = ['Smoking'],
evidence_card=[2])

model.add_cpds(cpd_B, cpd_R, cpd_C, cpd_S, cpd_F, cpd_M, cpd_Y)

[28] print(model.check_model())

True
```

```

from pgmpy.inference import VariableElimination
# Going to do variable elimination
infer = VariableElimination(model)

# Compute probability of Radiation given Cancer = 1
phi_query = infer.query(['Radiation'], evidence={'Cancer':1}, joint = False)
factor = phi_query['Radiation']
print('Probability of Radiation given Cancer = 1')
print(factor)

# Compute probability of Cancer given Skin Burn = 1
phi_query = infer.query(['Cancer'], evidence={'Skin Burn':1}, joint = False)
factor = phi_query['Cancer']
print('Probability of Cancer given Skin Burn = 1')
print(factor)

# Compute probability of Cancer given Using Microwave = 0
phi_query = infer.query(['Cancer'], evidence={'Using Microwave':0}, joint = False)
factor = phi_query['Cancer']
print('Probability of Cancer given Using Microwave = 0')
print(factor)

```

Probability of Radiation given Cancer = 1

Radiation	phi(Radiation)
Radiation(0)	0.6438
Radiation(1)	0.3562

Probability of Cancer given Skin Burn = 1

Cancer	phi(Cancer)
Cancer(0)	0.6092
Cancer(1)	0.3908

Probability of Cancer given Using Microwave = 0

Cancer	phi(Cancer)
Cancer(0)	0.8744
Cancer(1)	0.1256

2. $P(R = 1 \mid C = 1) = 0.6438$

3. $P(C = 1 \mid B = 1) = 0.3908$

4. Trails between Smoking and Using Microwave:

a. $S \rightarrow C \leftarrow R \leftarrow M$

Since Cancer is observed, the trail is active. Therefore Smoking and Using Microwave are independent.

5. $P(C = 1 \mid M = 0) = 0.1256$

Question 4

- Code and answers is submitted in zip file.

Question 5

a and b)

```
▼ Download MobileNetV2 model

▶ #<Write code for downloading MobileNetV2>

# Load the MobileNetV2 model pre-trained on ImageNet
base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(minSize, minSize, 3))

# Freeze the pre-trained layers
for layer in base_model.layers:
    layer.trainable = False
```

c)

✓ Add custom layers at the end of downloaded model

```
[ ] #<Write code for adding custom layers>
    from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
    from tensorflow.keras.models import Sequential
    # Create a new Sequential model
    model = Sequential()

    # Add the pre-trained base model
    model.add(base_model)

    # Add the rest of the layers
    model.add(GlobalAveragePooling2D())
    model.add(Dense(512, activation='relu'))
    model.add(Dense(10, activation='softmax')) # 10 classes for CIFAR-10
```

d)

▼ Add loss function, compile and train the model, and check accuracy on test data

```
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping

# Limit the training data and labels to 10,000 samples
limited_train_images = resized_train_images
limited_train_labels = train_labels[:10000]

# Compile the model with sparse categorical crossentropy
model.compile(optimizer=Adam(learning_rate=0.001), loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Define early stopping criteria
early_stopping = EarlyStopping(monitor='val_loss', patience=8, restore_best_weights=True)

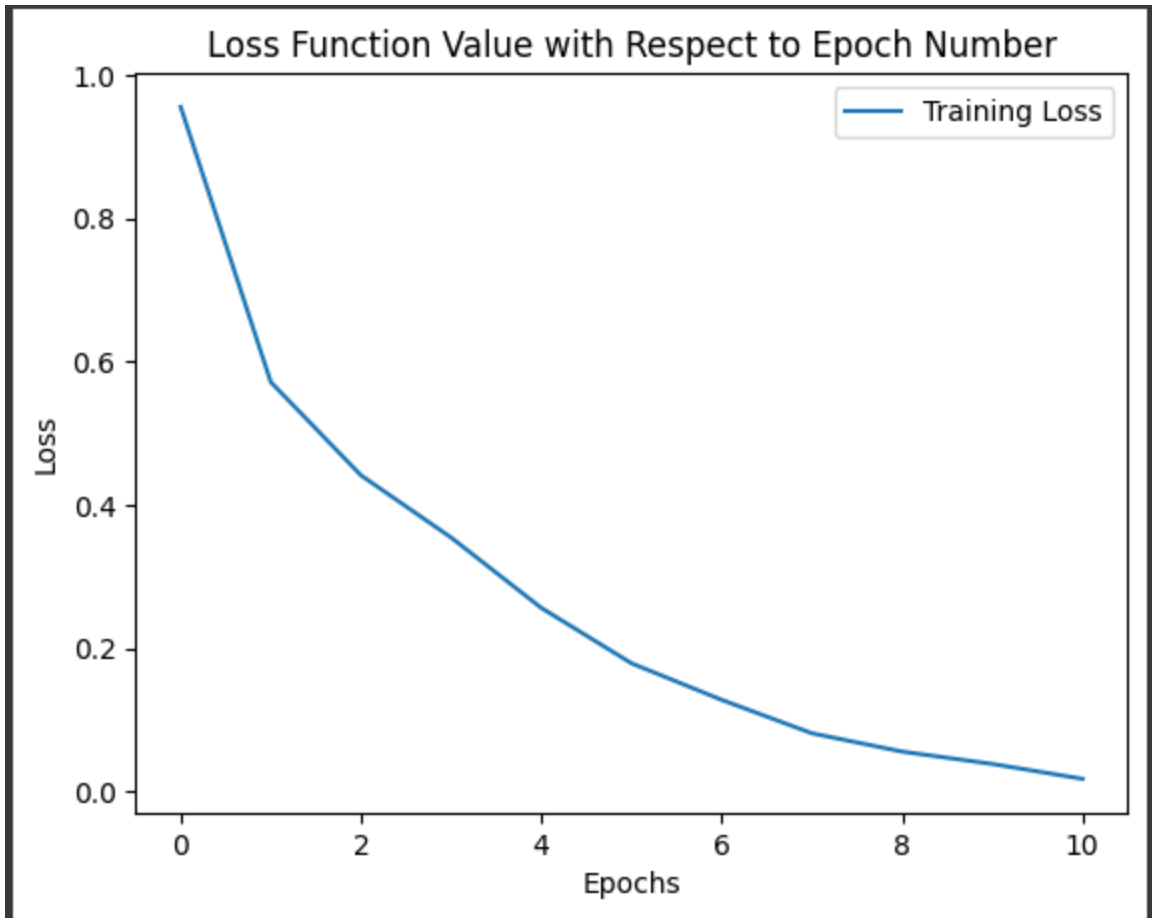
# Train the model with early stopping
history = model.fit(limited_train_images, limited_train_labels, batch_size=64, epochs=50,
                    validation_data=(resized_test_images, test_labels), callbacks=[early_stopping])

# Evaluate the model
test_loss, test_accuracy = model.evaluate(resized_test_images, test_labels, verbose=1)

print(f"Test accuracy: {test_accuracy * 100:.2f}%")
```

Epoch 1/50
157/157 [=====] - 20s 78ms/step - loss: 0.9554 - accuracy: 0.6790 - val_loss: 0.7351 - val_accuracy: 0.7427
Epoch 2/50
157/157 [=====] - 6s 38ms/step - loss: 0.5715 - accuracy: 0.7978 - val_loss: 0.7501 - val_accuracy: 0.7443
Epoch 3/50
157/157 [=====] - 6s 40ms/step - loss: 0.4413 - accuracy: 0.8450 - val_loss: 0.6973 - val_accuracy: 0.7705
Epoch 4/50
157/157 [=====] - 8s 53ms/step - loss: 0.3546 - accuracy: 0.8785 - val_loss: 0.7296 - val_accuracy: 0.7607
Epoch 5/50
157/157 [=====] - 9s 55ms/step - loss: 0.2564 - accuracy: 0.9164 - val_loss: 0.7614 - val_accuracy: 0.7579
Epoch 6/50
157/157 [=====] - 6s 40ms/step - loss: 0.1791 - accuracy: 0.9472 - val_loss: 0.7866 - val_accuracy: 0.7666
Epoch 7/50
157/157 [=====] - 6s 38ms/step - loss: 0.1281 - accuracy: 0.9646 - val_loss: 0.8054 - val_accuracy: 0.7637
Epoch 8/50
157/157 [=====] - 8s 54ms/step - loss: 0.0815 - accuracy: 0.9820 - val_loss: 0.8337 - val_accuracy: 0.7691
Epoch 9/50
157/157 [=====] - 8s 53ms/step - loss: 0.0560 - accuracy: 0.9904 - val_loss: 0.8981 - val_accuracy: 0.7639
Epoch 10/50
157/157 [=====] - 9s 56ms/step - loss: 0.0388 - accuracy: 0.9949 - val_loss: 0.9108 - val_accuracy: 0.7711
Epoch 11/50
157/157 [=====] - 9s 56ms/step - loss: 0.0180 - accuracy: 0.9993 - val_loss: 0.9346 - val_accuracy: 0.7729
313/313 [=====] - 4s 12ms/step - loss: 0.6973 - accuracy: 0.7705
Test accuracy: 77.05%

- 1) I extended the MobileNetV2 model by adding 1 GlobalAveragePooling2D layer, 1 dense layer with 512 nodes and ReLU activation and 1 dense output layer with 10 nodes and Softmax activation. Initially I added a MaxPool2D layer instead, but it was giving me very low training accuracy.
- 2)



- a. As seen from the screenshot for part (d), I added an EarlyStopping callback to the training process. This avoids overfitting and stops training as soon as `val_loss` stops improving. I set the patience parameter to 8 after experimenting with a few values and since the model reaches the required accuracy ($\geq 70\%$) in a few epochs, I set it to a relatively low number. I also set `restore_best_weights = True` so that when training is stopped, the model's weights will be rolled back to the state when it achieved the best performance on the validation loss.
- b. For the mini batch size, I experimented with a few values such as 32, 64 and 128 and decided on 64 as it gave me the highest training accuracy and a relatively smooth loss function with respect to epoch number graph.
- c. Similarly for learning rate, I experimented with different values ranging from 0.1 to 0.00001. I eventually settled on 0.001 as it gave me the necessary training accuracy when combined with the other set parameters.

3)

```
▼ Add loss function, compile and train the model, and check accuracy on test data

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping

# Limit the training data and labels to 10,000 samples
limited_train_images = resized_train_images
limited_train_labels = train_labels[:10000]

# Compile the model with sparse categorical crossentropy
model.compile(optimizer=Adam(learning_rate=0.001), loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Define early stopping criteria
early_stopping = EarlyStopping(monitor='val_loss', patience=8, restore_best_weights=True)

# Train the model with early stopping
history = model.fit(limited_train_images, limited_train_labels, batch_size=64, epochs=50,
                    validation_data=(resized_test_images, test_labels), callbacks=[early_stopping])

# Evaluate the model
test_loss, test_accuracy = model.evaluate(resized_test_images, test_labels, verbose=1)

print(f"Test accuracy: {test_accuracy * 100:.2f}%")

Epoch 1/50
157/157 [=====] - 20s 78ms/step - loss: 0.9554 - accuracy: 0.6790 - val_loss: 0.7351 - val_accuracy: 0.7427
Epoch 2/50
157/157 [=====] - 6s 38ms/step - loss: 0.5715 - accuracy: 0.7978 - val_loss: 0.7501 - val_accuracy: 0.7443
Epoch 3/50
157/157 [=====] - 6s 40ms/step - loss: 0.4413 - accuracy: 0.8450 - val_loss: 0.6973 - val_accuracy: 0.7705
Epoch 4/50
157/157 [=====] - 8s 53ms/step - loss: 0.3546 - accuracy: 0.8785 - val_loss: 0.7296 - val_accuracy: 0.7607
Epoch 5/50
157/157 [=====] - 9s 55ms/step - loss: 0.2564 - accuracy: 0.9164 - val_loss: 0.7614 - val_accuracy: 0.7579
Epoch 6/50
157/157 [=====] - 6s 40ms/step - loss: 0.1791 - accuracy: 0.9472 - val_loss: 0.7866 - val_accuracy: 0.7666
Epoch 7/50
157/157 [=====] - 6s 38ms/step - loss: 0.1281 - accuracy: 0.9646 - val_loss: 0.8054 - val_accuracy: 0.7637
Epoch 8/50
157/157 [=====] - 8s 54ms/step - loss: 0.0815 - accuracy: 0.9820 - val_loss: 0.8337 - val_accuracy: 0.7691
Epoch 9/50
157/157 [=====] - 8s 53ms/step - loss: 0.0560 - accuracy: 0.9904 - val_loss: 0.8981 - val_accuracy: 0.7639
Epoch 10/50
157/157 [=====] - 9s 56ms/step - loss: 0.0388 - accuracy: 0.9949 - val_loss: 0.9108 - val_accuracy: 0.7711
Epoch 11/50
157/157 [=====] - 9s 56ms/step - loss: 0.0180 - accuracy: 0.9993 - val_loss: 0.9346 - val_accuracy: 0.7729
313/313 [=====] - 4s 12ms/step - loss: 0.6973 - accuracy: 0.7705
Test accuracy: 77.05%
```

Note: I have added the Q3 code file in the zip folder as well in case the screenshot is not clear.